

Assignment 9

In the Simplesim assignment, we introduced a fictional machine called the Simplesim and the Simplesim Machine Language (SML). In this assignment, you are to write a compiler that converts programs written in a high-level programming language to the SML. This will require you to use the techniques that you learned to convert infix expressions to postfix and to evaluate those postfix expressions. We provide you with programs in this new high-level language with the intention of compiling them on the compiler that you will write and running the resulting SML program on the simulator you wrote.

1. The Simple Language

Before we begin building the compiler, we discuss a simple, yet powerful, high-level programming language called Simple.

- Every Simple statement consists of a line number and a Simple instruction. Line numbers must appear in ascending order.
- Each instruction begins with one of the following Simple commands: `ren`, `input`, `data`, `let`, `print`, `goto`, `if`, `...`, `goto`, or `end`.
- Simple evaluates only integer expressions using the arithmetic operators `+`, `-`, `*`, and `/`. The operators have the same precedence as in C++. Parentheses can be used to change the order of evaluation of an expression.

- All variables in a Simple program are lower case single letters. Simple uses only integer variables and constants. Simple does not have variable declarations – merely mentioning a variable name in a program causes the variable to be declared and initialized to zero automatically.
- Simple uses the conditional `if`...`goto` statement and the unconditional `goto` statement to alter the flow of control during program execution. If the condition in the `if`...`goto` statement is true, control is transferred to the specified line of the program. The following relational operators are valid in an `if`...`goto` statement: `<`, `>`, `<=`, `>=`, `=`, or `!=`. The precedence of these operators is the same as in C++.

Consider the following Simple program that accepts two input values and prints the maximum of the two:

```
10 ren      print the maximum of two numbers
11 input x
20 data 10
30 ren      get values
31 input x
32 ren      input y
60 ren      check x > y
61 ren      check x > y
62 ren      if x > y goto 111
70 ren      x is maximum, print y
80 ren      if x <= y goto 130
81 ren      y is maximum, print y
82 ren      if x <= y goto 130
100 goto 130
110 ren      x is maximum, print x
120 print x
130 data 20
990 end
```

Each line in a Simple program must start with a line number. The line numbers must appear in ascending order and must always start in the leftmost column, i.e., no leading whitespace. Execution of a Simple program starts at the first line. In our example program above, this is line 10, a `ren` command. `ren` commands are used to document the program and are not executed. Anything appearing after the word `ren` is simply ignored. Execution continues through the other documentation lines 11 and 12 and we reach line 20, a `data` command. `data` commands, like `ren` commands, are not executed. However, unlike `ren` commands, `data` commands do serve a purpose (explained below), `data` commands always end with a single integer. Execution proceeds through lines 20 and 30–32 and we approach line 40, an `input` command.

Input commands acquire a value from a data command and store that value in its variable. The very first `input` command that is executed uses the very first `data` command in the program. Subsequently executed input commands use the remaining `data` commands in the order in which they appear in the program. `data` commands are never used more than once. The same input command will consume many `data` commands if it executed many times (e.g., in a “loop”).

Line 40 is the first `input` command to be executed, so it acquires its value from the first `data` command (line 20) and stores the value 10 in the variable `x`. Line 50, our second input command to be executed, acquires its value from the next `data` command (line 130) and stores the value 20 in the variable `y`. Execution continues through the documentation lines 60–62 and we approach line 70, an `if`...`goto` command.

`if`...`goto` commands always of the form `if` <lop> <rellop> <exp> `goto` <linenum> where <lop> and <rellop> are always either a single variable or constant (i.e., not expressions). <rellop> is one of the relational operators listed above, and <linenum> is a line number appearing in the program.

Line 70 tests `x > y`. In our example program, this is false, so execution continues to line 80. Had the test been true, then execution would have continued at line 111. Note that line 111 is a `ren` command. Simple programs may branch (`if`...`goto` or `goto`) to `ren` and/or `data` lines, despite the fact that these lines are not executed. Execution continues through lines 80–82 and we approach line 90, a `print` command. `print` commands print either a single variable or constant. In our example program, the number 20 is printed and execution continues to line 100, a `goto` command, which branches (unconditionally) to line 130, our second `data` command. Since `data` commands are not executed, we proceed to line 990, an end command where program execution is terminated.

As a final example, consider the following Simple program that accepts a non-negative integer `x` and computes/prints the sum of integers from 1 through `x`.

```
10 ren sum 1 to x
15 input x
20 data 10
25 if x <= -1 goto 60
30 ren add x to total
50 goto 25
40 let x = x - 1
45 ren loop x
55 ren output result
60 print t
99 end
```

As always, execution starts with first line of the program. The first executable statement is line 15, an `input` command, which stores the value 10 (from the first `data` command on line 20) in variable `x`. Execution continues with line 25, an `if`...`goto` command where we test `x <= -1`. In our program this test is false, so execution continues to line 30 and add which add `x` to `t` and decrement `x`, respectively. `t` is the variable that we are using to accumulate the sum. The variable `t` has not been explicitly initialized, therefore the variables that appear in a Simple program are initialized (implicitly) to 0. Line 50 unconditionally branches back to line 25 where we test `x <= -1`. Eventually, that test will become true, in which case control will be transferred to line 60, the value of `t` (the sum) will be printed, and the program will terminate.

Simple does not provide a repetition structure (such as C++’s `for` or `while`). However, Simple can simulate each of C++’s repetition structures using the `if`...`goto` and `goto` commands as we did in the example above.

2. Input

The input to your program is a Simple program, just like the ones shown in the previous section. You may assume that all Simple programs given to your compiler are syntactically correct. Syntax checking is a critical component of any compiler, however, proper syntax checking is a very difficult problem and well beyond the scope of this course. Hence, when you write your compiler you may make the simplifying assumption that the programs you are compiling are syntactically correct. You may also assume that the programs you are compiling are semantically correct.

- All lines are numbered uniquely and appear in ascending order based on that line number.
- All line numbers start in the leftmost column.
- All line numbers appearing in `if`...`goto` and `goto` commands are in the program.
- There is at least one space between a line number and the command.
- There is at least one space between the command and the rest of the line.
- For `let` commands, there is at least one space before and after the `=` sign (note there may or may not be spaces in the expression on the right of the `=` sign).
- And that for `if`...`goto` commands, there will be at least one space separating the left and right operands from the relational operator.

3. Output

Your compiler will produce a complete SML program, including the -99999 line followed by any input (from any data in the Simple program) or an error message (error messages described in [Section 4](#)). The SML program generated by your compiler must be immediately suitable for execution on your Simplesim. In other words, you should be able to take the SML program generated by your compiler and feed it directly (without modification) as input to your `simplesim` program from Assignment 4.

4. The Two-Pass Compiler

The result of the compiler is the SML program, which is composed of SML instructions and data, a line containing -99999 to mark the end of the program, and possibly some input data. Your compiler program should use a memory array identical to the one it used in the Simplesim assignment and a data array with a counter to collect the values from the Simple `data` commands.

For example:

```
int memory[MEMSIZE];
int data[MEMSIZE];
int ndata;

// Your program will also use an array of flags and a symbol table (both are described below). The symbol table will have 1,000 "rows", where each row is an instance of the structure below. The flags array must have a flag for each memory location.
#define SYMBOL_TABLE_SIZE 1000
struct table_entry {
    int symbol;
    char type;           // 'C' constant, 'L' Simple line number, 'V' variable
    int location;        // Simplesim address (00 to MEMSIZE-1)
};

// The flags array will be an array of flags and a symbol table (both are described below). The symbol table will have 1,000 "rows", where each row is an instance of the structure below. The flags array must have a flag for each memory location.
table_entry symbol_table[SYMBOL_TABLE_SIZE];
int flags[MEMSIZE];
```

After some minor initialization, the compiler performs two passes. The first pass constructs a symbol table in which every line number, variable name, and constant of the Simple program is stored with its type and corresponding location in the final SML code (the symbol table is described in detail below).

Constants and variables are placed at the bottom of the Simplesim’s memory (memory locations 99, 98, 97, ...) and the SML program (instructions) are placed at the top of the Simplesim’s memory (locations 00, 01, 02, ...). The first pass also produces the corresponding SML instructions for each Simple statement. This results in a “hole” in the middle of the Simplesim’s memory (between the last SML instructions and the constants/variables) that will be used as stack space to evaluate the arithmetic expressions in the `let` commands.

As we will see, if the Simple program contains `let` statements or statements that transfer control to a line later in the program, the first pass results in an SML program containing some partial instructions and it uses the `flags` array to mark those instructions.

Your program should use the following variables to record the Simplesim address for the next instruction, the next constant or variable, and an index for the next symbol in the symbol table.

```
int next_instruction_addr;
int next_const_or_var_addr = MEMSIZE-1;
int next_symbol_table_idx;
```

The second pass of the compiler locates (using the `flags` array) and completes (using the symbol table and stack space) the partial instructions. After the second pass, the compiler prints the finished SML program, followed by any input data.

4.1. Initialization

Your compiler must initialize all of the Simplesim’s memory to 777 and all the flags to -1. It should then initialize the next instruction address to the “top” of the Simplesim’s memory (memory location 0), the next constant or variable address to the “bottom” of the Simplesim’s memory (memory location MEMSIZE-1), the index of the next entry into the symbol table to 0, and finally the counter for the Simple `data` commands to 0.

For example:

```
next_instruction_addr = 0;
next_const_or_var_addr = MEMSIZE-1;
next_symbol_table_idx = 0;
ndata = 0;
```

4.2. Pass 1

The first pass of the compiler reads and processes the Simple program one line at a time. In processing a single line of the Simple program, the compiler can make addition(s) to the symbol table, add (possibly partial) SML instruction(s) to the Simplesim’s memory, and add constants and/or variables to the Simplesim’s memory.

The general structure of the first pass should look like something like this:

```
string buffer1, buffer2, line_number, command;

while (getline(cin, buffer1))
{
    buffer2 = buffer1;           // buffer2 used for 'let'
    istringstream ss(buffer2);
    ss >> line_number;

    // ... code to add line_number to symbol table, type 'L' ...

    ss >> command;

    if (command == "input")
    {
        // ... code to process 'input' command ...
    }
    else if (command == "data")
    {
        // ... code to process 'data' command ...
    }
    else if (command == "let")
    {
        // ... code to process 'let' command ...
    }
    else if (command == "...")
    {
        // ... code to process 'rem' command ...
    }
    else if (command == "rem")
    {
        // ... code to process 'rem' command ...
    }
}
```

(The `istringstream` class is defined in the header file `<sstream>` and is part of the standard namespace.)

Each time a new line is read from the Simple program, its line number must be added to the symbol table along with the corresponding address of the next instruction of the SML program.

Each time a constant and/or variable is encountered, first the symbol table is checked to see if the constant or variable already exists in the symbol table.

- If the constant or variable does not exist within the symbol table (i.e., this is the first time that it has been encountered in the program) a Simplesim memory cell is allocated for it (initializing the memory to 0 if it was a variable) and the symbol table is updated by adding the variable or constant along with its memory location.
- If the variable or constant already exists in the symbol table, then no additional memory is allocated for it and the symbol table is left unmodified.

4.2.1. Processing Simple Commands

Here we discuss the details of how to process each of the Simple commands. In discussing each, it is assumed that the line number of the command has already been added to the symbol table.

- `ren`. There is nothing left to do after adding the line number to the symbol table.
- `data`. Add the integer appearing at the end of the line to the data array and increment `ndata`. If there is not enough room in the data array to add this element, print `*** ERROR: too many data lines ***` and terminate the compilation immediately, i.e., without printing the SML program.
- `input`. Search the symbol table for the variable appearing at the end of the line. If the variable appears in the symbol table, then extract the variable’s Simplesim memory location from the table. If the variable did not appear in the symbol table, allocate a Simplesim memory cell for it at `next_const_or_var_addr` and add it to the symbol table. Be sure to initialize the Simplesim’s memory for that variable to 0 and decrement `next_const_or_var_addr` for the next variable or constant. Be sure to note the newly-allocated memory location of the new variable (by saving it in the symbol table).
- Using the variable’s address, whether it was extracted from the symbol table or it was newly allocated, place a `READ` instruction at memory location `next_instruction_addr`. Be sure to increment `next_instruction_addr` for the next SML instruction.
- `print`. Process this in exactly the same manner as the `input` command, using the address from the symbol table if the variable already existed there or creating space for it and adding to the symbol table. The only difference in processing a `print` command is that you should place a `WRITE` instruction in memory (rather than a `READ`).
- `goto`. Search the symbol table for the line number appearing at the end of the line. If the line number appears in the symbol table (i.e., it refers to a line “before” this `goto` command), then extract the line number’s memory location from the symbol table and place/generate a `BRANCH` instruction at memory location `next_instruction_addr`. If the line number was not in the symbol table, then it refers to a line “after” this `goto` command. This is called a *forward* reference. In this case, we place/generate a partial `BRANCH` instruction (omitting the address to branch to) and flag the instruction to indicate that the second pass of the compiler must complete the instruction. Use the `flags` array and set `flags[next_instruction_addr]` to the line number this `goto` command is supposed to branch to. In either case, after writing the `BRANCH` instruction (partial or otherwise) be sure to increment `next_instruction_addr` for the next SML instruction.
- `if`...`goto`. Compilation of the `if`...`goto` and `let` statements is more complicated than the other statements – they are the only statements that produce more than one SML instruction. For an `if`...`goto` statement, the constants and/or variables appearing in the test are searched for in the symbol table and added to the Simplesim’s memory and symbol table if necessary (noting the Simplesim memory locations of each). Then the compiler produces SML instructions (typically a `LOAD` followed by a `SUBTRACT`) to test the condition and branch if appropriate. The result of the branch could be a forward reference and should be handled just like the forward references in the `goto` command i.e., writing partial `BRANCHNEG` and/or `BRANCHNEG` instructions and flagging them with the line number in the `flags` array. Each of the relational operators can be simulated using SML’s `BRANCHZERO` and `BRANCHNEG` instructions (or possibly a combination of both). Be sure to increment `next_instruction_addr` after each SML instruction.
- `let`. Compilation of the `let` statement is the most challenging, it not only produces more than one SML instruction but also the evaluation of the expression (rhs) requires you to manage your own stack (as an array) using the “hole” in the Simplesim’s memory between the SML instructions and the constants/variables. The problem is further complicated by the fact that during the first pass we do not know where that stack space starts, so we are forced to write *partial* SML instructions and flag them using the `flags` array.

The first thing you must do in processing a `let` statement is to search the symbol table for the variable being assigned. If it is not in the symbol table, then a Simplesim word must be allocated for it, initialized to 0, and it should be added to the symbol table. Then, using the copy of the statement (`buffer2`), call your `convert()` function from the previous assignment, passing it the address of the first non-whitespace character after the equal sign, to get a postfix expression. Evaluate the postfix expression using the stack algorithm to write the SML instructions (described in detail below).

The actual Simplesim memory location that will be the start of our stack cannot be known until after the space has been allocated for constants and variables (i.e., after the first pass has been completed). Once that is done, then we can use the “hole” between the last SML instruction and the last constant/variable. The stack will start above (but not immediately above) the last constant/variable added and grow “upward” toward the SML instructions each time we “push” a value. The first Simplesim word immediately above the last constant/variable is a special reserved word (the non-commutative operators - and /) (explained below), so the stack actually starts immediately above that specially reserved word.

Start the processing of the postfix expression by initializing a stack index `idx` to 0 (the next-stack-idx = 0), which represents the index into the stack where the next pushed value should go. Now you are ready to start processing the postfix expression.

Processing Operands

Each time you encounter an operand (constant or variable), you must search the symbol table for it, and if it is not there, you must add it to the Simplesim’s memory and the symbol table. Extract the operand’s memory location from the symbol table. Operands are pushed onto the stack. Since there are no memory-to-memory SML instructions, pushing a value onto the stack requires the use of the accumulator, i.e., you must `LOAD` and then `STORE`. Write the `LOAD` instruction using the operand’s address from the symbol table, then write a *partial* `STORE` instruction (omitting the address). Use the `flags` array to flag this `STORE` instruction by setting

```
flags[next_instruction_addr] = -3 - next_stack_idx;
```

Recall that the `flags` array was initialized with -1 and that it uses positive integers to represent forward referenced line numbers. This leaves negative numbers < -1 to represent stack indices. When the `flags` array element has the value -2, that refers to the specially reserved word for processing non-commutative operators (explained below). Values < -2 are used to represent stack indices, e.g., -3 means stack + 0, -4 means stack + 1, -5 means stack + 2, etc.

Thus, processing an operand requires finding its SML address from the symbol table (adding it if necessary), writing the `LOAD` instruction using that address, writing a partial `STORE` instruction, setting the `flags` array element corresponding to that `STORE` instruction to -3 - next_stack_idx, and incrementing next_stack_idx.

Processing Operators

Processing operators requires you to pop the two operands, perform the operation, and push the result.

Processing the commutative operators (+ and *) is a little easier than the non-commutative operators (- and /). Commutative operators pop the stack, placing the value into the accumulator (LOAD), and then perform the operation by popping the second operand. Then, using the copy of the statement (`buffer2`), call your `convert()` function from the previous assignment, passing it the address of the first non-whitespace character after the equal sign, to get a postfix expression. Evaluate the postfix expression using the stack algorithm to write the SML instructions (described in detail below).

The actual Simplesim memory location that will be the start of our stack cannot be known until after the space has been allocated for constants and variables (i.e., after the first pass has been completed). Once that is done, then we can use the “hole” between the last SML instruction and the last constant/variable. The stack will start above (but not immediately above) the last constant/variable added and grow “upward” toward the SML instructions each time we “push” a value. The first Simplesim word immediately above the last constant/variable is a special reserved word (the non-commutative operators - and /) (explained below), so the stack actually starts immediately above that specially reserved word.

Start the processing of the postfix expression by initializing a stack index `idx` to 0 (the next-stack-idx = 0), which represents the index into the stack where the next pushed value should go. Now you are ready to start processing the postfix expression.

Processing Operands

Each time you encounter an operand (constant or variable), you must search the symbol table for it, and if it is not there, you must add it to the Simplesim’s memory and the symbol table. Extract the operand’s memory location from the symbol table. Operands are pushed onto the stack. Since there are no memory-to-memory SML instructions, pushing a value onto the stack requires the use of the accumulator, i.e., you must `LOAD` and then `STORE`. Write the `LOAD` instruction using the operand’s address from the symbol table, then write a *partial* `STORE` instruction (omitting the address). Use the `flags` array to flag this `STORE` instruction by setting

```
flags[next_instruction_addr] = -3 - next_stack_idx;
```

Recall that the `flags` array was initialized with -1 and that it uses positive integers to represent forward referenced line numbers. This leaves negative numbers < -1 to represent stack indices. When the `flags` array element has the value -2, that refers to the specially reserved word for processing non-commutative operators (explained below). Values < -2 are used to represent stack indices, e.g., -3 means stack + 0, -4 means stack + 1, -5 means stack + 2, etc.

Thus, processing an operand requires finding its SML address from the symbol table (adding it if necessary), writing the `LOAD` instruction using that address, writing a partial `STORE` instruction, setting the `flags` array element corresponding to that `STORE` instruction to -3 - next_stack_idx, and incrementing next_stack_idx.

Processing Operators

Processing operators requires you to pop the two operands, perform the operation, and push the result.

Processing the commutative operators (+ and *) is a little easier than the non-commutative operators (- and /). Commutative operators pop the stack, placing the value into the accumulator (LOAD), and then perform the operation by popping the second operand. Then, using the copy of the statement (`buffer2`), call your `convert()` function from the previous assignment, passing it the address of the first non-whitespace character after the equal sign, to get a postfix expression. Evaluate the postfix expression using the stack algorithm to write the SML instructions (described in detail below).

For example, the code for processing the operator + would look like this:

```
memory[next_instruction_addr] = LOAD * 100;           // omit address
next_stack_idx--;
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = ADD * 100;           // for addition, omit address
next_stack_idx--;
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = STORE * 100;         // omit address
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = LOAD * 100;         // omit address
next_stack_idx--;
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = SUBTRACT * 100;      // for subtraction, omit address
flags[next_instruction_addr] = -3 - next_stack_idx;
next_stack_idx++;
next_instruction_addr++;

// Code for processing the operator + is nearly identical - just change ADD to MULTPLY.
```

Processing the non-commutative operators requires more care. Recall that the operand on the top of the stack is the *right* operand. Therefore, we cannot simply load that value into the accumulator and apply the operation using the next value on the stack. We must pop the right operand (LOAD) and temporarily store (STORE) it in the special memory location sitting just beneath the stack. Then pop the stack again, placing the left operand into the accumulator (LOAD). Apply the operator using the value in the special memory location (SUBTRACT or DIVIDE), leaving the result in the accumulator, and then push the result back onto the stack (STORE). Therefore, processing a non-commutative operator writes five *partial* SML instructions (omitting the address in all five).

For example, the code for processing the operator - would look like this:

```
memory[next_instruction_addr] = LOAD * 100;           // omit address
next_stack_idx--;
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = STORE * 100;         // omit address
flags[next_instruction_addr] = -2;
next_instruction_addr++;

memory[next_instruction_addr] = LOAD * 100;         // omit address
next_stack_idx--;
flags[next_instruction_addr] = -3 - next_stack_idx;
next_instruction_addr++;

memory[next_instruction_addr] = SUBTRACT * 100;      // for subtraction, omit address
flags[next_instruction_addr] = -3 - next_stack_idx;
next_stack_idx++;
next_instruction_addr++;

// Code for processing the operator - is nearly identical - just change SUBTRACT to DIVIDE.
```

Once you have finished evaluating the postfix expression, the answer is sitting on the top of the stack, where it must be removed and placed into the memory location of the variable of the `let` command. This is done with two SML instructions, a `partial` `LOAD` followed by a full `STORE`. The element of the `flags` array that corresponds to the partial `LOAD` instruction should be set to -3 (i.e., after the first instruction).

- end. Simply write a `HALT` instruction at memory location `next_instruction_addr`. Be sure to increment `next_instruction_addr` before writing the instruction.

There are only 100 words in the Simplesim memory and it is possible to deplete that. If by adding instructions you find that you have run past the end of memory or have entered the variable/constant section of the memory, or if by allocating space for variables/constants you have entered the program section of the memory, simply print the message `*** ERROR: ran out of Simplesim memory ***` and terminate the compilation immediately, i.e., without printing the SML program.

4.2.2. First Pass Example

To illustrate the first pass of the compiler, we take you through its steps (see [Figure 1](#) below) as it processes the summation example program from [Section 1](#).

We start with the first line of the program, line 10. As is done with each line of the program, the line number is added to the symbol table (see [Figure 3](#)), and since it is a `ren` command, nothing else is done.

Line 15 is an `input` command. The symbol table is searched for the variable `x` and it is not found (at this point there are only two entries in the symbol table, lines 10 and 15). Since `x` was not in the symbol table, a Simplesim word is allocated for `x` (at location 99), it is initialized to 0, and `x` is added to the symbol table. We are now ready to write our first SML instruction, a `READ` command using the address of `x` from the symbol table. Since we were able to write a full SML instruction (i.e., complete with address) we do not flag the instruction, meaning we leave this element of the `flags` array unchanged.

Compilation continues with line 20, a `data` command. The value (10) is placed into the first element of the data array and `ndata` is incremented.

Proceeding to line 25, we reach an `if`...`goto` command. The symbol table is searched for the left operand `x`, which is found. The symbol table is then searched for the right operand, the constant -1. It is not found in the symbol table, so the next available Simplesim word is allocated for -1 (at location 98), the constant -1 is added to the symbol table. We are now ready to start writing the SML instructions. This particular `if`...`goto` command is testing `<`. That test is assigned the variable `t`, `t` is initialized to 0, and `t` is added to the symbol table. We extract the Simplesim address (97) from the symbol table and save it for future use (the `STORE` operation at the end of processing this instruction). The infix expression `t < -1` is converted to its postfix equivalent `t x +` and compilation continues by processing the postfix expression, one symbol at a time.

The first symbol (`t`) is an operand, so we must push it onto the stack. This requires generating a `LOAD` and a `STORE`. The symbol table is searched for the symbol `t` and it is found. Its Simplesim address (99) is extracted from the symbol table (97) and used in the `LOAD` instruction. Now we must `STORE` the value in the stack. Since we do not know the Simplesim address where the stack will begin (we cannot know that until after the end of the first pass when space for all the variables/constants has been allocated), we are forced to write a *partial* `STORE` instruction (omitting the address) and flagging this instruction by placing a representation of the stack index (0) into the cell of the `flags` array corresponding to this instruction’s memory location (`flags[96] = -3 - next_stack_idx`) and the stack index is incremented. This concludes “pushing” the first operand (`t`) onto the stack and we continue with the next symbol in the postfix expression, `x`.

`x` is also an operand and must be pushed onto the stack. The symbol table is searched for the variable `x`. It too is found, and its Simplesim address (99) is extracted from the symbol table and is used in the `LOAD` instruction. Once again, we are forced to write a *partial* `STORE` instruction, this time using the next slot in the stack (stack index 1), of course incrementing the stack index after writing the `partial` `STORE` instruction and flagging its corresponding cell of the `flags` array. This concludes pushing the second operand (`x`) onto the stack.

Processing of the postfix expression continues as we move on to the next symbol, the commutative binary operator `+`. Processing such an operator requires us to pop the stack twice (performing the addition on the second pop) and pushing the result back onto the stack. We perform the first pop by decrementing the stack index and writing the *partial* `LOAD` instruction. We perform the second pop by decrementing the stack index and writing the *partial* `ADD` instruction. The sum is now sitting in the accumulator and must be pushed back onto the stack. This is performed by writing the *partial* `STORE` instruction and then incrementing the stack index.

This concludes the processing of the postfix expression. The stack should still have one value left, the value of the entire postfix expression. We must pop the stack and place the value into the variable that is the target of the `let` statement, in this case `t`. We pop the stack with a *partial* `LOAD` instruction and conclude with a full `STORE` instruction (using the address 97 that we extracted from the symbol table when we started processing this `let` command).

Compilation continues to line 40 where we encounter our second `let` instruction, this time involving the non-commutative binary operator `*`. We start again by searching the symbol table for the variable appearing to the left of the `=` sign (`x`) and find it, noting its Simplesim memory location (99). We continue by converting the infix expression `x - 1` to its postfix equivalent `x -` and - processing it one symbol at a time. The first two symbols of this postfix expression are also operands and they are both pushed onto the stack (SML instructions 14 + 27) the same way we pushed the first two operands in the previous `let` statement, with one minor difference: The constant -1 is not found in the symbol table, so Simplesim memory location 96 is allocated for it, set to -1, and the constant is added to the symbol table.

Now we must process the non-commutative binary operator `*`. Recall that the value on the top of the stack is the *right* operand. We cannot simply load that value into the accumulator and perform the subtraction operation using the second value on the stack. We must pop the right operand from the stack with a *partial* `LOAD` instruction and store it into the special memory location sitting just below the stack with a *partial* `STORE` instruction (specifying the special memory location in the `flags` array, `flags[100] = -2`). Processing continues by popping the next value from the stack into the accumulator and with a *partial* `LOAD` instruction and then writing the *partial* `SUBTRACT` instruction using the special memory location where we temporarily stored the right operand (`flags[21] = -2`). Finally, the difference is pushed back onto the stack with a *partial* `STORE` instruction. This concludes processing of the postfix expression and the result is popped off the stack and stored in `x` (SML instructions 28–95) in the same manner as the previous `let` command.

We continue through line 45 to line 50 where we encounter our first `goto` command, which always results in a `BRANCH` instruction. The only thing to determine is whether or not we can write a full instruction. We search the symbol table for the specified line number (25) and it is found. This allows us to write a full `BRANCH` instruction using the address that we extract from the symbol table. If the line number was not in the symbol table (e.g., a forward reference), we would have been forced to write a *partial* `BRANCH` instruction (omitting the address) and flag the instruction, setting `flags[25]` to the forward referenced line number.

Continuing to line 60, a `print` command, we search the symbol table for the specified variable (`t`). In this case, the variable was found. However, if the variable was not in the symbol table, space would have to be allocated for it and the symbol table would be updated. In either case, `print` statements always result in a full `WRITE` instruction. We conclude the first pass with line 99, an end command which generates a `HALT` instruction.

This concludes the compiler’s first pass through the summation program. At the end of the first pass, `data[0] = 10` and <